

# 1. Coding Design for the QnA System

## Components

### 1. Data Ingestion and Preprocessing:

- **Load and preprocess the dataset** to ensure it is in a format suitable for querying. This includes handling missing values, normalizing text, and possibly embedding the text data.

### 2. Retrieval-Augmented Generation (RAG) Model:

- **Retrieval:** Use embeddings and a search index to retrieve relevant documents.
- **Generation:** Use a generative model to formulate answers based on the retrieved documents.

### 3. Indexing:

- **Efficient Search:** Use vector databases like FAISS or Annoy for scalable and efficient similarity search.
- **Document Embeddings:** Convert text into embeddings to facilitate quick retrieval.

### 4. API Implementation:

- **Create an API** that allows querying the system and returns answers based on the dataset.

### 5. Scalability:

- **Load Balancing:** Use load balancers to distribute incoming requests across multiple instances.
- **Caching:** Cache frequently accessed data to reduce load times and improve response speed.
- **Horizontal Scaling:** Deploy the application on a cloud platform with auto-scaling capabilities.

```
import pandas as pd
```

```
import os
```

```
# Load the dataset
```

```
data_path =
```

```
'/content/drive/MyDrive/path/to/your/projects_with_embeddings.csv'
```

```
data = pd.read_csv(data_path)
```

```
# Preprocess the data

data.fillna("", inplace=True) data['text'] = data.apply(lambda row:
f"{row['project_name']} - {row['Unit type']} - Price: {row['price']}",
axis=1)

# Convert to a list of documents

documents = data.to_dict(orient='records')

texts = data['text'].tolist()
```

## **Creating an Index**

Use a vector database like FAISS for indexing and retrieving documents.

```
import faiss

import numpy as np from sentence_transformers

import SentenceTransformer # Initialize model

model = SentenceTransformer('all-MiniLM-L6-v2')

# Convert texts to embeddings

embeddings = model.encode(texts)

# Create FAISS index

dimension = embeddings.shape[1]

index = faiss.IndexFlatL2(dimension)

index.add(embeddings)
```

## **Building the API**

Use a web framework like FastAPI to build and expose the API.

```
from fastapi import FastAPI, HTTPException

from pydantic import BaseModel

import numpy as np
```

```

import faiss

app = FastAPI()

class Query(BaseModel):
    question: str

@app.post("/query")
def query_system(query: Query): # Convert query to embedding
    query_embedding = model.encode([query.question])

    # Retrieve documents

    distances, indices = index.search(query_embedding, k=5)

    retrieved_docs = [documents[idx] for idx in indices[0]]

    answer = " ".join([doc['text'] for doc in retrieved_docs])

    return {"answer": answer}

```

## Testing and Latency Reporting

### Functionality Testing:

1. **Test Endpoints:**
  - Ensure that API endpoints work correctly with various queries.
  - Use tools like Postman or `curl` to manually test.
2. **Automated Tests:**
  - Write test cases to validate different scenarios and edge cases.

### Performance Testing:

1. **Simulate Load:**
  - Use tools like `Apache JMeter` or `locust` to simulate multiple concurrent requests.
2. **Measure Latency:**
  - Record response times for different loads and configurations.

## Example

```
from locust import HttpUser, TaskSet, task, between

class UserBehavior(TaskSet):

    @task

    def query_api(self):

self.client.post("/query", json={"question": "What is the price of Project A?"})

class WebsiteUser(HttpUser):

tasks = [UserBehavior]

wait_time = between(1, 5)
```